

```

import tensorflow as tf
import tensorflow_hub as hub
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
from tensorflow.keras.utils import get_file
from sklearn.metrics import roc_curve, auc, confusion_matrix
from imblearn.metrics import sensitivity_score, specificity_score

import os
import glob
import zipfile
import random

# to get consistent results after multiple runs
tf.random.set_seed(7)
np.random.seed(7)
random.seed(7)

# 0 for benign, 1 for malignant
class_names = ["benign", "malignant"]

```

## Preparing the Dataset

```

def download_and_extract_dataset():
    # dataset from https://github.com/udacity/dermatologist-ai
    # 5.3GB
    train_url = "https://s3-us-west-1.amazonaws.com/udacity-
dlnd/datasets/skin-cancer/train.zip"
    # 824.5MB
    valid_url = "https://s3-us-west-1.amazonaws.com/udacity-
dlnd/datasets/skin-cancer/valid.zip"

```

```

# 5.1GB
test_url = "https://s3-us-west-1.amazonaws.com/udacity-
dlnd/datasets/skin-cancer/test.zip"
for i, download_link in enumerate([valid_url, train_url, test_url]):
    temp_file = f"temp{i}.zip"
    data_dir = get_file(origin=download_link,
fname=os.path.join(os.getcwd(), temp_file))
    print("Extracting", download_link)
    with zipfile.ZipFile(data_dir, "r") as z:
        z.extractall("data")
    # remove the temp file
    os.remove(temp_file)

# comment the below line if you already downloaded the dataset
download_and_extract_dataset()

```

```

# preparing data
# generate CSV metadata file to read img paths and labels from it
def generate_csv(folder, label2int):
    folder_name = os.path.basename(folder)
    labels = list(label2int)
    # generate CSV file
    df = pd.DataFrame(columns=["filepath", "label"])
    i = 0
    for label in labels:
        print("Reading", os.path.join(folder, label, "*"))
        for filepath in glob.glob(os.path.join(folder, label, "*")):
            df.loc[i] = [filepath, label2int[label]]
            i += 1
    output_file = f"{folder_name}.csv"
    print("Saving", output_file)
    df.to_csv(output_file)

# generate CSV files for all data portions, labeling nevus and
seborrheic keratosis

```

```
# as 0 (benign), and melanoma as 1 (malignant)
# you should replace "data" path to your extracted dataset path
# don't replace if you used download_and_extract_dataset() function
generate_csv("data/train", {"nevus": 0, "seborrheic_keratosis": 0,
"melanoma": 1})

generate_csv("data/valid", {"nevus": 0, "seborrheic_keratosis": 0,
"melanoma": 1})

generate_csv("data/test", {"nevus": 0, "seborrheic_keratosis": 0,
"melanoma": 1})
```

```
# loading data
train_metadata_filename = "train.csv"
valid_metadata_filename = "valid.csv"
# load CSV files as DataFrames
df_train = pd.read_csv(train_metadata_filename)
df_valid = pd.read_csv(valid_metadata_filename)
n_training_samples = len(df_train)
n_validation_samples = len(df_valid)
print("Number of training samples:", n_training_samples)
print("Number of validation samples:", n_validation_samples)
train_ds = tf.data.Dataset.from_tensor_slices((df_train["filepath"],
df_train["label"]))
valid_ds = tf.data.Dataset.from_tensor_slices((df_valid["filepath"],
df_valid["label"]))
```

Number of training samples: 2000  
Number of validation samples: 150

## Copy

Let's load the images:

```
# preprocess data
def decode_img(img):
    # convert the compressed string to a 3D uint8 tensor
    img = tf.image.decode_jpeg(img, channels=3)
    # Use `convert_image_dtype` to convert to floats in the [0,1] range.
    img = tf.image.convert_image_dtype(img, tf.float32)
    # resize the image to the desired size.
    return tf.image.resize(img, [299, 299])
```

```
def process_path(filepath, label):
    # load the raw data from the file as a string
    img = tf.io.read_file(filepath)
    img = decode_img(img)
    return img, label
```

```
valid_ds = valid_ds.map(process_path)
train_ds = train_ds.map(process_path)
# test_ds = test_ds
for image, label in train_ds.take(1):
    print("Image shape:", image.shape)
    print("Label:", label.numpy())
```

Image shape: (299, 299, 3)

Label: 0

```
# training parameters
```

```
batch_size = 64
```

```
optimizer = "rmsprop"
```

```
def prepare_for_training(ds, cache=True, batch_size=64,
    shuffle_buffer_size=1000):
```

```
    if cache:
```

```
        if isinstance(cache, str):
```

```
            ds = ds.cache(cache)
```

```
        else:
```

```
            ds = ds.cache()
```

```
    # shuffle the dataset
```

```
    ds = ds.shuffle(buffer_size=shuffle_buffer_size)
```

```
    # Repeat forever
```

```
    ds = ds.repeat()
```

```

# split to batches
ds = ds.batch(batch_size)
# `prefetch` lets the dataset fetch batches in the background while
the model
# is training.
ds = ds.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
return ds


valid_ds = prepare_for_training(valid_ds, batch_size=batch_size,
cache="valid-cached-data")
train_ds = prepare_for_training(train_ds, batch_size=batch_size,
cache="train-cached-data")
batch = next(iter(valid_ds))

def show_batch(batch):
    plt.figure(figsize=(12,12))
    for n in range(25):
        ax = plt.subplot(5,5,n+1)
        plt.imshow(batch[0][n])
        plt.title(class_names[batch[1][n].numpy()].title())
        plt.axis('off')

show_batch(batch)

```

Output:



```
# building the model
# InceptionV3 model & pre-trained weights
module_url = "https://tfhub.dev/google/tf2-
preview/inception_v3/feature_vector/4"
m = tf.keras.Sequential([
    hub.KerasLayer(module_url, output_shape=[2048], trainable=False),
    tf.keras.layers.Dense(1, activation="sigmoid")
])

m.build([None, 299, 299, 3])
m.compile(loss="binary_crossentropy", optimizer=optimizer,
metrics=["accuracy"])
m.summary()
```

Model: "sequential"

| Layer (type)                     | Output Shape | Param #  |
|----------------------------------|--------------|----------|
| =====                            |              |          |
| keras_layer (KerasLayer)         | multiple     | 21802784 |
| =====                            |              |          |
| dense (Dense)                    | multiple     | 2049     |
| =====                            |              |          |
| Total params: 21,804,833         |              |          |
| Trainable params: 2,049          |              |          |
| Non-trainable params: 21,802,784 |              |          |
| =====                            |              |          |

## Training the Model

We now have our dataset and the model, let's get them together:

```
model_name = f"benign-vs-malignant_{batch_size}_{optimizer}"
tensorboard =
tf.keras.callbacks.TensorBoard(log_dir=os.path.join("logs",
model_name))
# saves model checkpoint whenever we reach better weights
modelcheckpoint = tf.keras.callbacks.ModelCheckpoint(model_name +
"_{val_loss:.3f}.h5", save_best_only=True, verbose=1)

history = m.fit(train_ds, validation_data=valid_ds,
                steps_per_epoch=n_training_samples // batch_size,
                validation_steps=n_validation_samples // batch_size,
verbose=1, epochs=100,
                callbacks=[tensorboard, modelcheckpoint])
```

Here is a part of the output during training:

```
Train for 31 steps, validate for 2 steps
Epoch 1/100
30/31 [=====>.] - ETA: 9s - loss: 0.4609 -
accuracy: 0.7760
```

```

Epoch 00001: val_loss improved from inf to 0.49703, saving model to
benign-vs-malignant_64_rmsprop_0.497.h5
31/31 [=====] - 282s 9s/step - loss: 0.4646 -
accuracy: 0.7722 - val_loss: 0.4970 - val_accuracy: 0.8125
<..SNIPED..>
Epoch 27/100
30/31 [=====>.] - ETA: 0s - loss: 0.2982 -
accuracy: 0.8708
Epoch 00027: val_loss improved from 0.40253 to 0.38991, saving model to
benign-vs-malignant_64_rmsprop_0.390.h5
31/31 [=====] - 21s 691ms/step - loss: 0.3025
- accuracy: 0.8684 - val_loss: 0.3899 - val_accuracy: 0.8359
<..SNIPED..>
Epoch 41/100
30/31 [=====>.] - ETA: 0s - loss: 0.2800 -
accuracy: 0.8802
Epoch 00041: val_loss did not improve from 0.38991
31/31 [=====] - 21s 690ms/step - loss: 0.2829
- accuracy: 0.8790 - val_loss: 0.3948 - val_accuracy: 0.8281
Epoch 42/100
30/31 [=====>.] - ETA: 0s - loss: 0.2680 -
accuracy: 0.8859
Epoch 00042: val_loss did not improve from 0.38991
31/31 [=====] - 21s 693ms/step - loss: 0.2722
- accuracy: 0.8831 - val_loss: 0.4572 - val_accuracy: 0.8047

```

## Model Evaluation

First, let's load our test set, just like previously:

```

# evaluation
# load testing set
test_metadata_filename = "test.csv"
df_test = pd.read_csv(test_metadata_filename)
n_testing_samples = len(df_test)
print("Number of testing samples:", n_testing_samples)
test_ds = tf.data.Dataset.from_tensor_slices((df_test["filepath"],
df_test["label"]))

def prepare_for_testing(ds, cache=True, shuffle_buffer_size=1000):

```



```

if cache:
    if isinstance(cache, str):
        ds = ds.cache(cache)
    else:
        ds = ds.cache()
ds = ds.shuffle(buffer_size=shuffle_buffer_size)
return ds

test_ds = test_ds.map(process_path)
test_ds = prepare_for_testing(test_ds, cache="test-cached-data")

```

The above code loads our test data and prepares it for testing:

Number of testing samples: 600

600 images of the shape (299, 299, 3) can fit our memory, let's convert our test set from `tf.data` into a NumPy array:

```

# convert testing set to numpy array to fit in memory (don't do that
when testing
# set is too large)
y_test = np.zeros((n_testing_samples,))
X_test = np.zeros((n_testing_samples, 299, 299, 3))
for i, (img, label) in enumerate(test_ds.take(n_testing_samples)):
    # print(img.shape, label.shape)
    X_test[i] = img
    y_test[i] = label.numpy()

print("y_test.shape:", y_test.shape)
# load the weights with the least loss
m.load_weights("benign-vs-malignant_64_rmsprop_0.390.h5")
print("Evaluating the model...")
loss, accuracy = m.evaluate(X_test, y_test, verbose=0)
print("Loss:", loss, " Accuracy:", accuracy)

```

Copy

Output:

Evaluating the model...

Loss: 0.4476394319534302    Accuracy: 0.8

The below function does that:

```
def get_predictions(threshold=None):
    """
    Returns predictions for binary classification given `threshold`
    For instance, if threshold is 0.3, then it'll output 1 (malignant)
    for that sample if
    the probability of 1 is 30% or more (instead of 50%)
    """
    y_pred = m.predict(X_test)
    if not threshold:
        threshold = 0.5
    result = np.zeros((n_testing_samples,))
    for i in range(n_testing_samples):
        # test melanoma probability
        if y_pred[i][0] >= threshold:
            result[i] = 1
        # else, it's 0 (benign)
    return result

threshold = 0.23
# get predictions with 23% threshold
# which means if the model is 23% sure or more that is malignant,
# it's assigned as malignant, otherwise it's benign
y_pred = get_predictions(threshold)
```

Now let's draw our confusion matrix and interpret it:

```
def plot_confusion_matrix(y_test, y_pred):
    cmn = confusion_matrix(y_test, y_pred)
    # Normalise
    cmn = cmn.astype('float') / cmn.sum(axis=1)[:, np.newaxis]
```


```

# print it
print(cmn)
fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(cmn, annot=True, fmt='.2f',
            xticklabels=[f"pred_{c}" for c in class_names],
            yticklabels=[f"true_{c}" for c in class_names],
            cmap="Blues"
            )
plt.ylabel('Actual')
plt.xlabel('Predicted')
# plot the resulting confusion matrix
plt.show()

plot_confusion_matrix(y_test, y_pred)

```

Output:



```


def plot_roc_auc(y_true, y_pred):
    """
    This function plots the ROC curves and provides the scores.
    """
    # prepare for figure
    plt.figure()
    fpr, tpr, _ = roc_curve(y_true, y_pred)
    # obtain ROC AUC
    roc_auc = auc(fpr, tpr)
    # print score
    print(f"ROC AUC: {roc_auc:.3f}")
    # plot ROC curve
    plt.plot(fpr, tpr, color="blue", lw=2,
             label='ROC curve (area = {f:.2f})'.format(d=1,
f=roc_auc))
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC curves')
    plt.legend(loc="lower right")
    plt.show()

plot_roc_auc(y_test, y_pred)

```

[Copy](#)

Output:



ROC AUC : 0.671